Deep Learning Final Project

May 12, 2025

Predicting Temperature Anomalies

My project was to predict future temperature anomalies—i.e. the global temperature difference from the average for a day.

Predicting Temperature Anomalies

My project was to predict future temperature anomalies—i.e. the global temperature difference from the average for a day.In essence, this is a time series modelling problem where we have 1-dimensional data and we want to predict what comes next in the series.

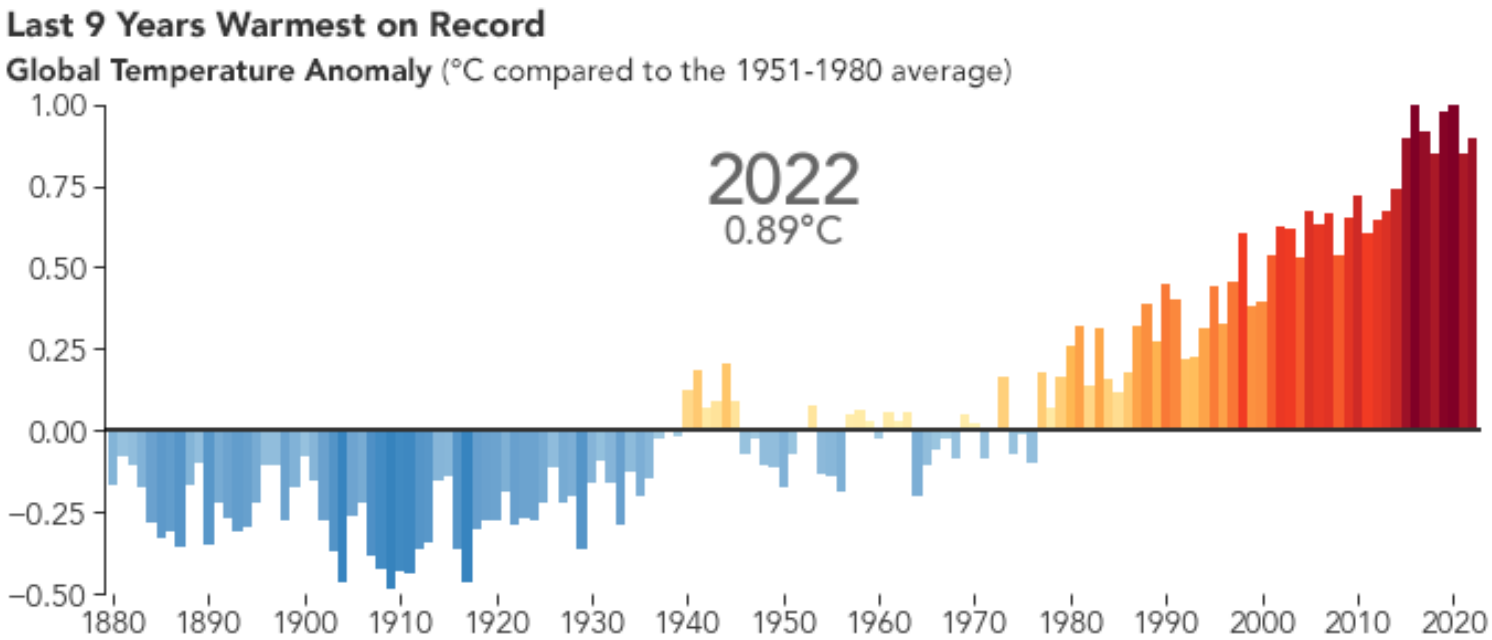

How to Do This

In order to do this, I needed a dataset. UC Berkeley has temperature anomaly information ranging from 1880 to 2022. This was used to make the graphic you saw on the last slide.

How to Do This

In order to do this, I needed a dataset. UC Berkeley has temperature anomaly information ranging from 1880 to 2022. This was used to make the graphic you saw on the last slide.

From there, we need a model. Because this is a time series problem, we want a network which cares about the order of our elements.

How to Do This

In order to do this, I needed a dataset. UC Berkeley has temperature anomaly information ranging from 1880 to 2022. This was used to make the graphic you saw on the last slide.

From there, we need a model. Because this is a time series problem, we want a network which cares about the order of our elements. I started with an LSTM model which became too cumbersome. File sizes were too large and there was a tendency to overfit.

How to Do This

In order to do this, I needed a dataset. UC Berkeley has temperature anomaly information ranging from 1880 to 2022. This was used to make the graphic you saw on the last slide.

From there, we need a model. Because this is a time series problem, we want a network which cares about the order of our elements. I started with an LSTM model which became too cumbersome. File sizes were too large and there was a tendency to overfit.

I settled on a temporal convolutional network. These took less time to train, had smaller file sizes and were less prone to overfitting.

Results

The results were pretty good! Overfitting was much lower on a reasonably-sized network, network sizes were kept down and they were relatively quick to train.

Results

The results were pretty good! Overfitting was much lower on a reasonably-sized network, network sizes were kept down and they were relatively quick to train.Here, the red line is the prediction and the blue line is our actual data set.

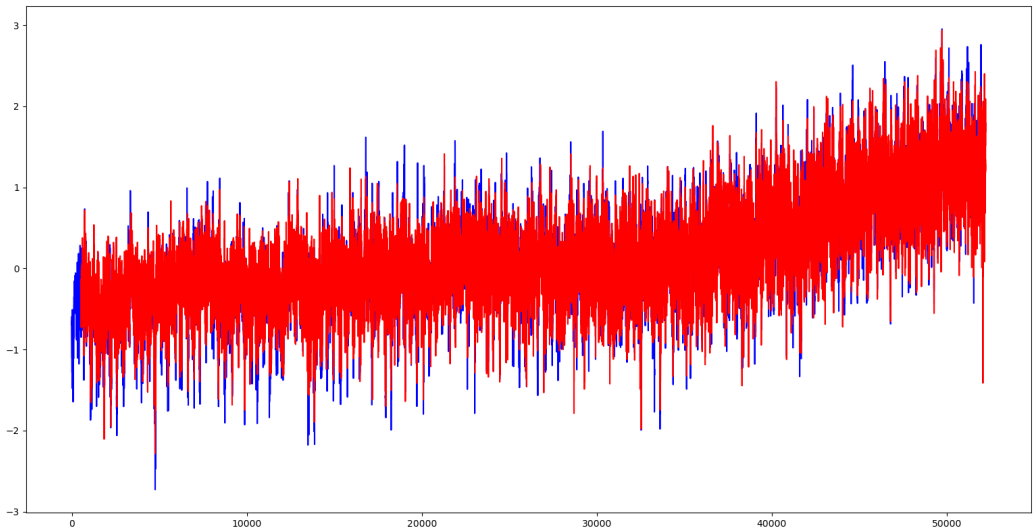

Figure 1: Predictions v. actual data

Results (cont.)

Unfortunately, my networks were quite horrific at actually forecasting.

Results (cont.)

Unfortunately, my networks were quite horrific at actually forecasting.